

Lecture 7

The Network Layer

Data Plane


Subjects of today:

- The Role of the Network Layer
- The Router
- Internet Protocol Addressing
- Beyond IPv4, NAT and IPv6
- Generalized forwarding in brief

7.1 The Role of the Network Layer


The Internet Protocol Stack

- ***application***: supporting network applications
 - FTP, SMTP, HTTP
- ***transport***: process-process data transfer
 - TCP, UDP
- ***network***: routing of datagrams from source to destination
 - IP, routing protocols
- ***link***: data transfer between neighboring network elements
 - Ethernet, 802.111 (WiFi), PPP
- ***physical***: bits “on the wire”




Introduction 1-60

Layered protocol stack




Every layer must...




- offer services to an upper layer.
- comply with agreed protocols.
- utilize services from the underlying layer.

The Network Layer at the host...



- Retrieves and delivers segments from/to the Transport Layer.
- Comply with network layer protocols:
 - Data plane: IP (IPv4, IPv6).
 - Control plane: E.g. ICMP.
- Sends and receives datagrams to/from other layer 3 devices through the Link Layer.


The Network Layer at the router...



- Comply with network layer protocols:
 - Data plane: IP (IPv4, IPv6).
 - Control plane: E.g. ICMP.
- Sends and receives datagrams to/from other layer 3 devices through the Link Layer.
- Forward datagrams from one interface (physical port) to another.
- Route datagrams to the “right” interface.

Network Layer

- transport segment from sending to receiving host
- on sending side encapsulates segments into datagrams
- on receiving side, delivers segments to transport layer
- network layer protocols in *every* host & router
- router examines header fields in all IP datagrams passing through it




Two Key Network-Layer Functions

network-layer functions:

- *forwarding*: move packets from router's input to appropriate router output
- *routing*: determine route taken by packets from source to destination
 - *routing algorithms*

analogy: taking a trip


- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



Network layer: Data plane vs. Control plane

Data plane

- local, per-router function
- determines how datagram arriving on router input port is forwarded to router output port
- forwarding function




Control plane

- network-wide logic
- determines how datagram is routed among routers along end-to-end path from source host to destination host
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers


Per-router Control Plane

Individual routing algorithm components *in each and every router* interact in the control plane



Logically Centralized Control Plane

A distinct (typically remote) controller interacts with local control agents (CAs)




Network Service Model

Possibilities

- Guaranteed delivery (loss)
- Guaranteed delivery with bounded delay (time)
- In-order packet delivery
- Minimal bandwidth
- Security


Actual

- Best-effort service...




7.2 The Router

The Router




Router Architecture Overview

- high-level view of generic router architecture:



Input port functions




physical layer:
bit-level reception

data link layer:
e.g., Ethernet
see chapter 6

decentralized switching:

- using **header field** values, **lookup** output port using forwarding table in input port memory ("*match plus action*")
- **goal**: complete input port processing at 'line speed'
- queuing: if datagrams arrive faster than forwarding rate into switch fabric

Input port functions



- using header field values, lookup output port using forwarding table in input port memory ("*match plus action*")
- ***destination-based forwarding***: forward based only on destination IP address (traditional)
- ***generalized forwarding***: forward based on any set of header field values

Destination-based forwarding

forwarding table

| Destination Address Range | Link Interface |
|---|----------------|
| 11001000 00010111 00010 000 00000000 through 11001000 00010111 00010 111 11111111 | 0 |
| 11001000 00010111 00011000 00000000 through 11001000 00010111 00011000 11111111 | 1 |
| 11001000 00010111 00011001 00000000 through 11001000 00010111 00011111 11111111 | 2 |
| otherwise | 3 |

Q: but what happens if ranges don't divide up so nicely?

Longest prefix matching

longest prefix matching —

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

| Destination Address Range | Link interface |
|----------------------------------|----------------|
| 11001000 00010111 00010*** ***** | 0 |
| 11001000 00010111 00011000 ***** | 1 |
| 11001000 00010111 00011*** ***** | 2 |
| otherwise | 3 |

examples:

DA: 11001000 00010111 00010110 10100001



which interface?

DA: 11001000 00010111 00011000 10101010

which interface?

Switching fabrics


- transfer packet from input buffer to appropriate output buffer
- switching rate: rate at which packets can be transferred from inputs to outputs
- three types of switching fabrics:



What are the drawbacks of each method?

Input port queuing


- fabric slower than input ports combined -> queueing may occur at input queues
 - *queueing delay and loss due to input buffer overflow!*
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward



output port contention:
only one red datagram can be
transferred.
lower red packet is blocked

one packet time later:
green packet
experiences HOL
blocking

Output ports




- *buffering* required when datagrams arrive from fabric faster than the transmission rate
- *scheduling discipline* chooses among queued datagrams for transmission

Datagram (packets) can be lost due to congestion/lack of buffers

Priority scheduling – who gets best performance?

Scheduling Mechanisms


- **Scheduling:** choose next packet to send on link
- **FIFO (first in first out) scheduling:** send in order of arrival to queue
 - **Discard policy:** if packet arrives to full queue: who to discard?
 - Tail drop: drop arriving packet
 - Priority: drop/remove on priority basis
 - Random: drop/remove randomly




7.3 Internet Protocol Addressing

The Internet Network Layer

Host and router network layer functions:




IP datagram format



IP fragmentation, reassembly

- network links have MTU (max.transfer size) - largest possible link-level frame
 - different link types, different MTUs
- large IP datagram divided ("fragmented") within net
 - one datagram becomes several datagrams
 - "reassembled" only at final destination
 - IP header bits used to identify, order related fragments



IP Fragmentation & Reassembly

example:




- ❖ 4000 byte datagram
- ❖ MTU = 1500 bytes

1480 bytes in data field

offset =
 $1480/8$




one large datagram becomes several smaller datagrams



What happens if two hosts sends with the same ID?

IP Addressing: Introduction

- *IP address:* 32-bit identifier for host, router *interface*
- *interface:* connection between host/router and physical link
 - router's typically have multiple interfaces
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- *IP addresses associated with each interface*



223.1.1.1 = $\underline{11011111} \underline{00000001} \underline{00000001} \underline{00000001}$

223 1 1 1


IP Addressing: Introduction

*Q: how are interfaces
actually connected?*

*A: we'll learn about that
in chapter 6,7*

*A: wired Ethernet interfaces
connected by Ethernet switches*


For now: don't need to worry about
how one interface is connected to
another (with no intervening router)



*A: wireless WiFi interfaces
connected by WiFi base station*

Subnets

- IP address:
 - subnet part - high order bits
 - host part - low order bits
- *What's a subnet ?*
 - device interfaces with same subnet part of IP address
 - can physically reach each other *without intervening router*




network consisting of 3 subnets

Subnets

recipe

- to determine the subnets, detach each interface from its host or router, creating islands of isolated networks
- each isolated network is called a *subnet*



subnet mask: /24

IP addressing: CIDR


CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- address format: $a.b.c.d/x$, where x is # bits in subnet portion of address



Subnets again...

How many subnets if the subnet mask was: /16?



IP addresses: How to get one?

- Hard-coded by system admin in a file
- DHCP: **D**ynamic **H**ost **C**onfiguration **P**rotocol: dynamically get address from as server
 - “plug-and-play”

DHCP: Dynamic Host Configuration Protocol


Goal: allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network

DHCP overview:

- host broadcasts “**DHCP discover**” msg [optional]
- DHCP server responds with “**DHCP offer**” msg [optional]
- host requests IP address: “**DHCP request**” msg
- DHCP server sends address: “**DHCP ack**” msg

DHCP client-server scenario



DHCP: more than IP addresses

DHCP can return more than just an allocated IP address on subnet:

- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)

IP addresses: how to get one?

Q: How does a *network* get the subnet part of its IP address?

A: Gets the allocated portion of its provider ISP's address space

| | | |
|----------------|---|----------------|
| ISP's block | <u>11001000</u> <u>00010111</u> <u>00010000</u> <u>00000000</u> | 200.23.16.0/20 |
| Organization 0 | <u>11001000</u> <u>00010111</u> <u>00010000</u> <u>00000000</u> | 200.23.16.0/23 |
| Organization 1 | <u>11001000</u> <u>00010111</u> <u>00010010</u> <u>00000000</u> | 200.23.18.0/23 |
| Organization 2 | <u>11001000</u> <u>00010111</u> <u>00010100</u> <u>00000000</u> | 200.23.20.0/23 |
| ... | | |
| Organization 7 | <u>11001000</u> <u>00010111</u> <u>00011110</u> <u>00000000</u> | 200.23.30.0/23 |

IP addressing: the last word...


Q: How does an ISP get block of addresses?

A: ICANN: Internet Corporation for Assigned Names and Numbers <https://www.icann.org/>

- allocates addresses
- manages DNS
- assigns domain names, resolves disputes

7.4 Beyond IPv4, NAT and IPv6

NAT: Network Address Translation



all datagrams *leaving* local network have *same* single source NAT IP address: $138.76.29.7$, different source port numbers

datagrams with source or destination in this network have $10.0.0/24$ address for source, destination (as usual)

NAT: Network Address Translation

Motivation: Local network uses just one IP address as far as outside world is concerned:

- Range of addresses not needed from ISP: just one IP address for all devices
- Can change addresses of devices in local network without notifying outside world
- Can change ISP without changing addresses of devices in local network
- Devices inside local net not explicitly addressable, visible by outside world (a security plus)

NAT: Network Address Translation


- *implementation:* NAT router must:
 - *outgoing datagrams:* replace (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
 - . . . remote clients/servers will respond using (NAT IP address, new port #) as destination addr
 - *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
 - *incoming datagrams:* replace (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: Network Address Translation

2: NAT router changes datagram source addr from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table

| NAT translation table | |
|-----------------------|----------------|
| WAN side addr | LAN side addr |
| 138.76.29.7, 5001 | 10.0.0.1, 3345 |
| | |

1: host 10.0.0.1 sends datagram to 128.119.40.186, 80




3: reply arrives dest. address: 138.76.29.7, 5001

4: NAT router changes datagram dest addr from 138.76.29.7, 5001 to 10.0.0.1, 3345

NAT: Network Address Translation

- 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
 - Routers should only process up to layer 3
 - Address shortage should be solved by IPv6
 - Violates end-to-end argument
 - NAT possibility must be taken into account by app designers, e.g., P2P applications



What if a client
wants to
connect to
server behind
NAT?

Port Forwarding

| Quick Select | |
|--------------------|---------------|
| Famous Server List | Please select |
| Famous Game List | Please select |

| Custom Configuration | |
|----------------------|-----------------|
| Service Name | NAS * Optional |
| Protocol | BOTH |
| External Port | 5200 |
| Internal Port | 5100 * Optional |
| Internal IP Address | 192.168.1.10 |
| Source IP | * Optional |

* External Port
The External Port accepts the following formats
1. Port ranges using a colon ":" between the starting and ending port, such as 300:350.
2. Single ports using a comma "," between individual ports, such as 566, 789.
3. A Mix of port ranges and single ports, using colons ":" and commas ",", such as 1015:1024, 3021.

* Source IP
If you want to open your port to a specific IP address from the internet, input the IP address you want to specify in the Source IP field.

Cancel OK

IPv6: Motivation

- 32-bit address space completely allocated¹
- Additional motivation:
 - header format helps speed processing/forwarding
 - header changes to facilitate QoS

IPv6 datagram format:


- fixed-length 40 byte header
 - no fragmentation allowed
 - no checksum
 - no options
- *ICMPv6:* new version of ICMP
 - additional message types, e.g. “Packet Too Big”
 - multicast group management functions

IPv6 Datagram Format

priority: identify priority among datagrams in flow


flow Label: identify datagrams in same “flow.” (concept of “flow” not well defined).

next header: identify upper layer protocol for data (in the payload)




Transitioning from IPv4 to IPv6

- Not all routers can be upgraded simultaneously
 - no “flag days”
 - how will network operate with mixed IPv4 and IPv6 routers?
- *Tunneling:* IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers



Tunneling




IPv6: adoption

- Google: 30% of clients access services via IPv6
- NIST: 1/3 of all US government domains are IPv6 capable
- *Long (long!) time for deployment, use*
 - 25 years and counting!
 - think of application-level changes in last 25 years: WWW, social media, streaming media, gaming, telepresence, ...
 - *Why?*

7.5 Generalized forwarding in brief


Per-router Control plane

Individual routing algorithm components *in each and every router* interact in the control plane




Generalized Forwarding and SDN

Each router contains a *flow table* that is computed and distributed by a *logically centralized* routing controller



OpenFlow data plane abstraction


- *Flow*: defined by header fields
- Generalized forwarding: simple packet-handling rules
 - *Pattern*: match values in packet **header fields**
 - *Actions*: *for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
 - *Priority*: disambiguate overlapping patterns
 - *Counters*: #bytes and #packets



Flow table in a router (computed and distributed by controller) define router's match+action rules

OpenFlow data plane abstraction


- *flow*: defined by header fields
- generalized forwarding: simple packet-handling rules
 - *Pattern*: match values in packet header fields
 - *Actions: for matched packet*: drop, forward, modify, matched packet or send matched packet to controller
 - *Priority*: disambiguate overlapping patterns
 - *Counters*: #bytes and #packets



* : wildcard

1. $\text{src}=1.2.*.*$, $\text{dest}=3.4.5.* \rightarrow \text{drop}$
2. $\text{src} = *.*.*.*$, $\text{dest}=3.4.*.* \rightarrow \text{forward}(2)$
3. $\text{src}=10.1.2.3$, $\text{dest} = *.*.*.* \rightarrow \text{send to controller}$

OpenFlow: Flow Table Entries



Next session, we will talk about the control plane of the network layer.

Read chapter 5 in the book (page 407-467).

Today we introduce the second hand-in assignment of this course, you must hand in a journal of your work with Wireshark Lab. 6 in the hand-in folder at itsLearning. Deadline is 14/12-2025.

As with the first hand-in "Python Lab 3" the hand-in is mandatory for all that adhere to the "Data communications" T590000101 course description.

The conditions of the hand in, can be found in the "Lab 6" plan at itslearning/resources.